

Overview

This notebook is created for solving the [Kaggle House Price Modeling Challenge](#). The dataset has various categorical (nominal, binary, and ordinal variables) and numerical variables to determine the selling price of a house. The training dataset is rather small (1460) and it has a lot of variables (81). The lack of data and the huge number of variables was a pretty big challenge in this competition.

This notebook is my first base model to solving the problem. I perform a bit of EDA (exploratory data analysis) and I cycle through several models such as Linear Regression all the way to Xgboost and I also use a bit of model stacking to optimize model performance.

Table of Contents

1. Utility Functions
2. Handling Missing Values
3. Understanding Target Variable
4. Modeling

In [27]:

```
import sys

sys.path.append('C:\\\\Users\\\\Nikhil\\\\Data_Science_Projects')
from common_ds_modules import missing_values, data_manipulation, modeling
import os
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import Lasso
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import RandomizedSearchCV
from sklearn.tree import DecisionTreeRegressor
from xgboost import XGBRegressor
import pickle
from sklearn.linear_model import Lasso, Ridge, ElasticNet
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import StackingRegressor
```

Utility Functions

```
In [28]: def get_column(df, query=''):  
    for c in df.columns:  
        if query.lower() in c.lower():  
            print(f'Column: {c}')
```

```
In [29]: MAX_MODELS = 100
```

```
In [30]: train_df = pd.read_csv('train.csv')
        test_df = pd.read_csv('test.csv')
```

```
In [31]: train_df['MSSubClass'] = train_df['MSSubClass'].apply(lambda x: str(x))
test_df['MSSubClass'] = test_df['MSSubClass'].apply(lambda x: str(x))
train_df['LogSalePrice'] = np.log(train_df['SalePrice'] + 1)
```

Fill in Missing Values

```
In [33]: rename_columns = {'1stFlrSF': 'FirstFlrSF', '2ndFlrSF':'SecondFlrSF', '3SsnPorch':'Thr  
  
train_df = train_df.rename(rename_columns, axis='columns')  
test_df = test_df.rename(rename_columns, axis='columns')
```

Getting rid of all columns that have more than 30% missing values. I chose 30 as a rather arbitrary number, I thought it was just right, if it is more than 30, then you will corrupt the variable distribution with the fill in value.

For categorical variables I fill in missing values by using the mode, and for numerical variables I use the mean.

```
In [34]: train_drop = missing_values.get_high_missing_value_columns(train_df, 30)
test_drop = missing_values.get_high_missing_value_columns(test_df, 30)

train_low_missing_vals_df = train_df.drop(train_drop, axis='columns')
test_low_missing_vals_df = test_df.drop(test_drop, axis='columns')
```

This function gets the numerical and categorical variables, and fills in missing values as well.

```
In [35]: numerical_variables, categorical_variables = data_manipulation.get_numerical_categorical()
```

Sanity check to make sure that training and testing dataset don't have any columns with missing values after fill in.

```
In [36]: missing_values.get_null_columns(train_low_missing_vals_df)
missing_values.get_null_columns(test_low_missing_vals_df)
```

```
In [37]: def plot_variable_dist(df, variables):
    for c in variables:
        plt.hist(train_df[c])
        plt.title(f'Distribution for {c}')
        plt.show()
```

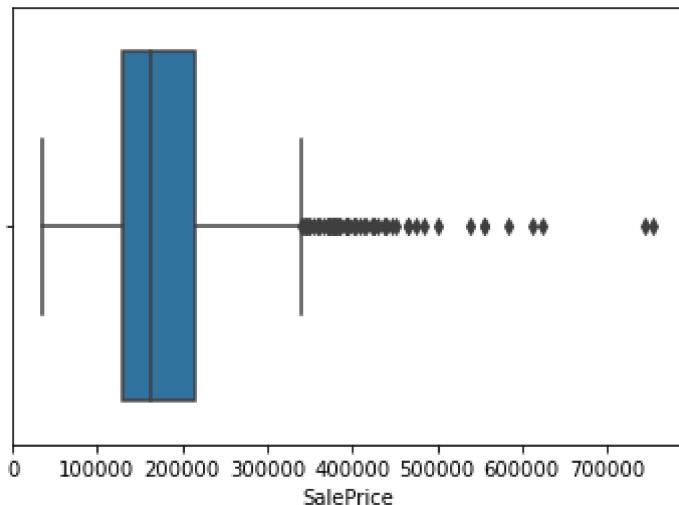
Understanding the target variable

For this problem I am trying to predict the SalePrice, so that is the target. In this section I will create a few plots to examine the SalePrice variable

```
In [38]: target = 'SalePrice'
```

```
In [39]: sns.boxplot(train_df[target])
plt.show()
```

C:\Users\Nikhil\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning:
Pass the following variable as a keyword arg: x. From version 0.12, the only valid po
sitional argument will be `data`, and passing other arguments without an explicit key
word will result in an error or misinterpretation.
warnings.warn(



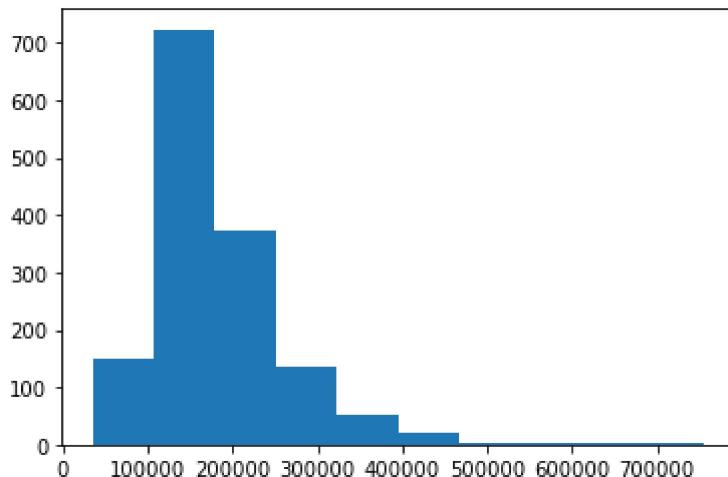
```
In [40]: q1 = np.percentile(train_df[target], 25)
q3 = np.percentile(train_df[target], 75)
iqr = q3 - q1
max_outlier = q3+1.5*iqr
print(f'Max outlier: {max_outlier}')
num_outliers = train_df[train_df['SalePrice'] > max_outlier].shape[0]
print(f'Percentage of outliers: {100*round(num_outliers/train_df.shape[0], 2)}')
```

Max outlier: 340037.5
Percentage of outliers: 4.0

Histogram of Sale Price

```
In [41]: plt.hist(train_df['SalePrice'])
```

```
Out[41]: (array([148., 723., 373., 135., 51., 19., 4., 3., 2., 2.]),
array([ 34900., 106910., 178920., 250930., 322940., 394950., 466960.,
538970., 610980., 682990., 755000.]),
<BarContainer object of 10 artists>)
```



```
In [42]: label = 'SalePrice'
```

Correlation Matrix for all numerical variables

```
In [43]: corr = train_df.corr()[label].sort_values()
valid_correlations = corr[(corr > 0.4) & (corr < 1)].index.tolist()
train_df[valid_correlations].corr()
```

	Fireplaces	MasVnrArea	GarageYrBlt	YearRemodAdd	YearBuilt	TotRmsAbvGrd	FullBath
Fireplaces	1.000000	0.249070	0.046822	0.112581	0.147716	0.326114	0.243671
MasVnrArea	0.249070	1.000000	0.252691	0.179618	0.315707	0.280682	0.276105
GarageYrBlt	0.046822	0.252691	1.000000	0.642277	0.825667	0.148112	0.484121
YearRemodAdd	0.112581	0.179618	0.642277	1.000000	0.592855	0.191740	0.439151
YearBuilt	0.147716	0.315707	0.825667	0.592855	1.000000	0.095589	0.468271
TotRmsAbvGrd	0.326114	0.280682	0.148112	0.191740	0.095589	1.000000	0.554121
FullBath	0.243671	0.276833	0.484557	0.439046	0.468271	0.554784	1.000000
FirstFlrSF	0.410531	0.344501	0.233449	0.240379	0.281986	0.409516	0.380105
TotalBsmtSF	0.339519	0.363936	0.322445	0.291066	0.391452	0.285573	0.323105
GarageArea	0.269141	0.373066	0.564567	0.371600	0.478954	0.337822	0.405105
GarageCars	0.300789	0.364204	0.588920	0.420622	0.537850	0.362289	0.469105
GrLivArea	0.461679	0.390857	0.231197	0.287389	0.199010	0.825489	0.630105
OverallQual	0.396765	0.411876	0.547766	0.550684	0.572323	0.427452	0.550105
LogSalePrice	0.489450	0.430809	0.541073	0.565608	0.586570	0.534422	0.594105

```
In [44]: # compartmentalize in another program
```

```
In [45]: y_train = train_low_missing_vals_df['LogSalePrice']
train_model_data, test_model_data = data_manipulation.get_train_test_model_data(train_
                                         numeri
                                         'SaleP
train_drop_columns = list(set(train_model_data.columns) - set(test_model_data.columns))
train_model_data_final = train_model_data.drop(train_drop_columns, axis='columns')
```

Modeling

Default random forest regressor

```
In [46]: import pickle
```

```
In [47]: %%time
rf_model_default = RandomForestRegressor().fit(train_model_data_final, y_train)
modeling.evaluate(train_model_data_final, y_train, rf_model_default)
filename = 'rf_default_model.pkl'
pickle.dump(rf_model_default, open(filename, 'wb'))
```

mean RMSLE = 0.14209999066690118, std RMSLE = 0.00972142378293112
CPU times: total: 24 s
Wall time: 33.7 s

```
In [48]: #modeling section in new program
```

RandomForest Regressor

```
In [49]: %%time
rf_param_grid = {
    'n_estimators': [30, 40],
    'max_depth':[2,3],
    'min_samples_split':[20, 30],
    'min_samples_leaf':[20, 30]
}

rf_model = modeling.get_random_forest_regressor_results(train_model_data_final, y_train)

model = RandomForestRegressor(**rf_model.best_params_).fit(train_model_data_final, y_train)
modeling.evaluate(train_model_data_final, y_train, model)
```

C:\Users\Nikhil\anaconda3\lib\site-packages\sklearn\model_selection_search.py:306: UserWarning: The total space of parameters 16 is smaller than n_iter=100. Running 16 iterations. For exhaustive searches, use GridSearchCV.

```
warnings.warn(
mean RMSLE = 0.19647765698202765, std RMSLE = 0.006551179403259499
CPU times: total: 35.2 s
Wall time: 44.5 s
```

Decision Tree

In [50]:

```
%time

dt_param_grid = {
    'max_depth':[2,3],
    'min_samples_split':[50, 30],
    'min_samples_leaf':[50, 30],
    'max_features' : [30, 40],
    'min_impurity_decrease': [i/100 for i in range(1, 3)],
    'ccp_alpha': [i/10 for i in range(1, 3)]
}

dt_model = modeling.get_decision_tree_model(train_model_data_final, y_train, dt_param_
modeling.evaluate(train_model_data_final, y_train, dt_model)
```

C:\Users\Nikhil\anaconda3\lib\site-packages\sklearn\model_selection_search.py:306: UserWarning: The total space of parameters 64 is smaller than n_iter=100. Running 64 iterations. For exhaustive searches, use GridSearchCV.
 warnings.warn(
 C:\Users\Nikhil\anaconda3\lib\site-packages\sklearn\model_selection_search.py:306: UserWarning: The total space of parameters 64 is smaller than n_iter=100. Running 64 iterations. For exhaustive searches, use GridSearchCV.
 warnings.warn(
 C:\Users\Nikhil\anaconda3\lib\site-packages\sklearn\model_selection_search.py:306: UserWarning: The total space of parameters 64 is smaller than n_iter=100. Running 64 iterations. For exhaustive searches, use GridSearchCV.
 warnings.warn(
 C:\Users\Nikhil\anaconda3\lib\site-packages\sklearn\model_selection_search.py:306: UserWarning: The total space of parameters 64 is smaller than n_iter=100. Running 64 iterations. For exhaustive searches, use GridSearchCV.
 warnings.warn(
 C:\Users\Nikhil\anaconda3\lib\site-packages\sklearn\model_selection_search.py:306: UserWarning: The total space of parameters 64 is smaller than n_iter=100. Running 64 iterations. For exhaustive searches, use GridSearchCV.
 warnings.warn(
 C:\Users\Nikhil\anaconda3\lib\site-packages\sklearn\model_selection_search.py:306: UserWarning: The total space of parameters 64 is smaller than n_iter=100. Running 64 iterations. For exhaustive searches, use GridSearchCV.
 warnings.warn(
 mean RMSLE = 0.3992256860383698, std RMSLE = 0.015929916505866743
 CPU times: total: 38.9 s
 Wall time: 50.6 s

XGBRegressor

In [51]:

```
%time

xgb_param_grid = {
    'n_estimators': [30, 40],
    'max_depth':[2,3],
    'max_leaves':[20, 30],
    'reg_alpha': [i/100 for i in range(1,3)],
    'reg_lambda': [i/100 for i in range(1,3)],
    'colsample_bytree': [i/10 for i in range(1, 3)],
    'min_child_weight': [10*i for i in range(1,3)],
    'learning_rate': [i/100 for i in range(1,3)]
}

xgb_model = modeling.get_xgb_regressor_model(train_model_data_final, y_train, xgb_param
```

```

modeling.evaluate(train_model_data_final, y_train, xgb_model)

xgb_regressor = XGBRegressor(**xgb_model.best_params_)
scores = cross_val_score(xgb_regressor, train_model_data_final, y_train, scoring='neg_
print(f'mean RMSLE = {-np.mean(scores)}', std RMSLE = {np.std(scores)})')

from sklearn.metrics import mean_squared_error, mean_squared_log_error, mean_absolute_
xgb_regressor = XGBRegressor(**xgb_model.best_params_).fit(train_model_data_final, y_t
train_pred = xgb_regressor.predict(train_model_data_final)
train_pred = np.exp(train_pred) - 1
print(f'MAE = {mean_absolute_error(train_df["SalePrice"], train_pred)}')
print(f'RMSLE = {np.sqrt(mean_squared_error(train_df["SalePrice"], train_pred))}')
test_pred = xgb_regressor.predict(test_model_data.drop('MSSubClass_150', axis='columns
test_pred = np.exp(test_pred) - 1
test_df['PredictedSalePrice'] = test_pred
test_df[['Id', 'PredictedSalePrice']].to_csv('xgb_predicted.csv', index=False)

```

mean RMSLE = 5.142351050249056, std RMSLE = 0.01570371335182782
mean RMSLE = 5.1422521013551625, std RMSLE = 0.01544287904286843
MAE = 179928.26673997854
RMSLE = 196612.05255440506
CPU times: total: 1h 4min 4s
Wall time: 12min 37s

Linear Regression

In [52]:

```

lr_model = modeling.get_linear_regressor_model(train_model_data_final, y_train)
modeling.evaluate(train_model_data_final, y_train, lr_model)

mean RMSLE = 0.17140003259393077, std RMSLE = 0.027991473177328228

```

Lasso Regression

In [53]:

```

%%time
lasso_param_grid = {
    'alpha':[i/100 for i in range(1, 3)]
}
lasso_gs = modeling.get_lasso_regression_model(train_model_data_final, y_train, lasso_
lasso = Lasso(**lasso_gs.best_params_)
scaler = StandardScaler().fit(train_model_data_final, y_train)
scores = cross_val_score(lasso, scaler.transform(train_model_data_final), y_train, sco
print(f'mean RMSLE = {-np.mean(scores)}')

mean RMSLE = 0.14418176386928863
CPU times: total: 656 ms
Wall time: 408 ms

```

Elastic Net Regression

In [54]:

```

%%time
elastic_net_param_grid = {
    'alpha': [100*i for i in range(1,11)] + [i/100 for i in range(1,3)],
    'l1_ratio':[i/100 for i in range(1,3)]
}

elastic_net_model = modeling.get_elastic_net_regressor_model(train_model_data_final, )

```

```
best_elastic_net_model = ElasticNet(**elastic_net_model.best_params_)
modeling.evaluate(StandardScaler().fit_transform(train_model_data_final), y_train, best_elastic_net_model)

mean RMSLE = 0.15132605569444632, std RMSLE = 0.02434875878994365
CPU times: total: 7.61 s
Wall time: 3.16 s
```

In [55]: `elastic_net_model.best_params_`

Out[55]: `{'alpha': 0.02, 'l1_ratio': 0.02}`

In [56]: `#el_net = ElasticNet(**elastic_net_model.best_params_)
scores = cross_val_score(best_elastic_net_model, StandardScaler().fit_transform(train_model_data_final), y_train, scoring='neg_mean_squared_error')
print(f'mean RMSLE = {-np.mean(scores)}')`

mean RMSLE = 0.15132605569444632

KNN Regressor

In [57]: `%time`

```
knn_param_grid = {
    'weights' : ['uniform', 'distance'],
    'algorithm' : ['brute', 'ball_tree', 'kd_tree'],
    'n_neighbors' : [10*i for i in range(1,3)],
    'leaf_size': [10*i for i in range(1,3)],
    'p': [i for i in range(1,3)],
}

knn_grid_search = modeling.get_knn_regressor_model(train_model_data_final, y_train, knn_param_grid)
knn_regressor = KNeighborsRegressor(**knn_grid_search.best_params_)
modeling.evaluate(train_model_data_final, y_train, knn_regressor)
```

C:\Users\Nikhil\anaconda3\lib\site-packages\sklearn\model_selection_search.py:306: UserWarning: The total space of parameters 48 is smaller than n_iter=100. Running 48 iterations. For exhaustive searches, use GridSearchCV.

```
warnings.warn(
mean RMSLE = 0.20376538868577163, std RMSLE = 0.006544415044060678
CPU times: total: 1min 8s
Wall time: 47.6 s
```

Stacking Regressor with Lasso, Ridge, Elastic, Random Forest, and XGBRegressor as the final estimator

In [58]: `estimators = [
 ('KNN', KNeighborsRegressor(**knn_grid_search.best_params_)),
 ('Lasso', Lasso(**lasso_gs.best_params_)),
 ('Ridge', Ridge(alpha=580)),
 ('Elastic', ElasticNet(**elastic_net_model.best_params_)),
 ('RandomForestRegressor', RandomForestRegressor())
]

final_estimator=XGBRegressor(**xgb_model.best_params_)

stacking_regressor = modeling.get_stacking_regressor(train_model_data_final, y_train, estimators)
modeling.evaluate(train_model_data_final, y_train, stacking_regressor)`

```
C:\Users\Nikhil\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.147e+00, tolerance: 1.509e-02
    model = cd_fast.enet_coordinate_descent()
mean RMSLE = 5.141168118210303, std RMSLE = 0.011603195447551994
```

In [59]:

```
%%time
stacking_regressor = stacking_regressor.fit(train_model_data_final, y_train)
test_pred = stacking_regressor.predict(test_model_data.drop('MSSubClass_150', axis='co
test_pred = np.exp(test_pred) - 1
test_df['PredictedSalePrice'] = test_pred
test_df[['Id', 'PredictedSalePrice']].to_csv('stacking_stacking_regressor_xgb_pr
CPU times: total: 30.6 s
Wall time: 35.9 s
```

In [60]:

```
%%time
estimators = [
    ('KNN', KNeighborsRegressor(**knn_grid_search.best_params_)),
    ('Lasso', Lasso(**lasso_gs.best_params_)),
    ('Ridge', Ridge(alpha=580)),
    ('Elastic', ElasticNet(**elastic_net_model.best_params_)),
    ('XGBRegressor', XGBRegressor(**xgb_model.best_params_))
]

final_estimator=RandomForestRegressor()

stacking_regressor = modeling.get_stacking_regressor(train_model_data_final, y_train,
modeling.evaluate(train_model_data_final, y_train, stacking_regressor)
```

```
C:\Users\Nikhil\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.147e+00, tolerance: 1.509e-02
    model = cd_fast.enet_coordinate_descent()
mean RMSLE = 0.1300674332791923, std RMSLE = 0.010709599685574183
CPU times: total: 1min 24s
Wall time: 22 s
```

In [61]:

```
%%time
stacking_regressor = stacking_regressor.fit(train_model_data_final, y_train)
test_pred = stacking_regressor.predict(test_model_data.drop('MSSubClass_150', axis='co
test_pred = np.exp(test_pred) - 1
test_df['PredictedSalePrice'] = test_pred
test_df[['Id', 'PredictedSalePrice']].to_csv('stacking_regressor_rf.csv', index=False)

CPU times: total: 17.3 s
Wall time: 5.17 s
```

Both stacking regressors have far lower RMSLE's than the other models, it just goes to show you how powerful stacking is