

SerieA_ML

July 24, 2022

1 Serie A Machine Learnin Project

1.1 Importing Data

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib
import numpy as np
%matplotlib inline
```

Data from: <https://fbref.com/en/comps/11/Serie-A-Stats>

```
[2]: table_df = pd.read_csv('seriea_table.csv')
```

```
[3]: table_df
```

```
[3]:
```

	Rk	Squad	MP	W	D	L	GF	GA	GD	Pts	Pts/MP	xG	xGA	\
0	1	Milan	38	26	8	4	69	31	38	86	2.26	63.1	34.8	
1	2	Inter	38	25	9	4	84	32	52	84	2.21	81.4	39.2	
2	3	Napoli	38	24	7	7	74	31	43	79	2.08	59.4	31.9	
3	4	Juventus	38	20	10	8	57	37	20	70	1.84	51.5	38.0	
4	5	Lazio	38	18	10	10	77	58	19	64	1.68	55.8	48.3	
5	6	Roma	38	18	9	11	59	43	16	63	1.66	63.7	38.5	
6	7	Fiorentina	38	19	5	14	59	51	8	62	1.63	58.8	44.1	
7	8	Atalanta	38	16	11	11	65	48	17	59	1.55	66.2	45.2	
8	9	Hellas Verona	38	14	11	13	65	59	6	53	1.39	51.8	48.9	
9	10	Torino	38	13	11	14	46	41	5	50	1.32	48.8	39.9	
10	11	Sassuolo	38	13	11	14	64	66	-2	50	1.32	56.0	67.0	
11	12	Udinese	38	11	14	13	61	58	3	47	1.24	53.1	52.0	
12	13	Bologna	38	12	10	16	44	55	-11	46	1.21	44.1	54.2	
13	14	Empoli	38	10	11	17	50	70	-20	41	1.08	46.6	67.8	
14	15	Sampdoria	38	10	6	22	46	63	-17	36	0.95	37.0	58.5	
15	16	Spezia	38	10	6	22	41	71	-30	36	0.95	39.3	67.6	
16	17	Salernitana	38	7	10	21	33	78	-45	31	0.82	37.7	65.9	
17	18	Cagliari	38	6	12	20	34	68	-34	30	0.79	39.5	61.5	
18	19	Genoa	38	4	16	18	27	60	-33	28	0.74	37.8	51.6	
19	20	Venezia	38	6	9	23	34	69	-35	27	0.71	36.1	72.6	

	xGD	xGD/90	Attendance	Top Team Scorer \
0	28.3	0.74	44015	Olivier Giroud Rafael Leão - 11
1	42.2	1.11	44473	Lautaro Martínez - 21
2	27.6	0.73	28119	Victor Osimhen - 14
3	13.4	0.35	22621	Paulo Dybala - 10
4	7.6	0.20	23263	Ciro Immobile - 27
5	25.2	0.66	41929	Tammy Abraham - 17
6	14.7	0.39	21107	Dušan Vlahović - 17
7	20.9	0.55	10447	Mario Pašalić - 13
8	2.9	0.08	13894	Giovanni Simeone - 17
9	8.9	0.23	9846	Andrea Belotti - 8
10	-10.9	-0.29	8362	Gianluca Scamacca - 16
11	1.0	0.03	12144	Gerard Deulofeu - 13
12	-10.1	-0.27	14158	Marko Arnautović - 14
13	-21.2	-0.56	6356	Andrea Pinamonti - 13
14	-21.5	-0.57	9417	Francesco Caputo - 11
15	-28.3	-0.74	6709	Daniele Verde - 8
16	-28.2	-0.74	15073	Federico Bonazzoli - 10
17	-22.0	-0.58	9718	João Pedro - 13
18	-13.8	-0.36	12326	Mattia Destro - 9
19	-36.5	-0.96	6648	Thomas Henry - 9

	Goalkeeper	Notes
0	Mike Maignan	→ Champions League via league finish
1	Samir Handanović	→ Champions League via league finish
2	David Ospina	→ Champions League via league finish
3	Wojciech Szczęsny	→ Champions League via league finish
4	Thomas Strakosha	→ Europa League via league finish
5	Rui Patrício	→ Europa League via league finish
6	Pietro Terracciano	→ Europa Conference League via league finish
7	Juan Musso	NaN
8	Lorenzo Montipò	NaN
9	Vanja Milinković-Savić	NaN
10	Andrea Consigli	NaN
11	Marco Silvestri	NaN
12	Łukasz Skorupski	NaN
13	Guglielmo Vicario	NaN
14	Emil Audero	NaN
15	Ivan Provedel	NaN
16	Vid Belec	NaN
17	Alessio Cragno	Relegated
18	Salvatore Sirigu	Relegated
19	Niki Mäenpää	Relegated

```
[4]: table_df.drop(columns=['Notes', 'Goalkeeper', 'Top Team Scorer', 'MP'],
    ↪inplace=True)
```

```
[5]: table_df
```

```
[5]:
```

	Rk	Squad	W	D	L	GF	GA	GD	Pts	Pts/MP	xG	xGA	xGD	\
0	1	Milan	26	8	4	69	31	38	86	2.26	63.1	34.8	28.3	
1	2	Inter	25	9	4	84	32	52	84	2.21	81.4	39.2	42.2	
2	3	Napoli	24	7	7	74	31	43	79	2.08	59.4	31.9	27.6	
3	4	Juventus	20	10	8	57	37	20	70	1.84	51.5	38.0	13.4	
4	5	Lazio	18	10	10	77	58	19	64	1.68	55.8	48.3	7.6	
5	6	Roma	18	9	11	59	43	16	63	1.66	63.7	38.5	25.2	
6	7	Fiorentina	19	5	14	59	51	8	62	1.63	58.8	44.1	14.7	
7	8	Atalanta	16	11	11	65	48	17	59	1.55	66.2	45.2	20.9	
8	9	Hellas Verona	14	11	13	65	59	6	53	1.39	51.8	48.9	2.9	
9	10	Torino	13	11	14	46	41	5	50	1.32	48.8	39.9	8.9	
10	11	Sassuolo	13	11	14	64	66	-2	50	1.32	56.0	67.0	-10.9	
11	12	Udinese	11	14	13	61	58	3	47	1.24	53.1	52.0	1.0	
12	13	Bologna	12	10	16	44	55	-11	46	1.21	44.1	54.2	-10.1	
13	14	Empoli	10	11	17	50	70	-20	41	1.08	46.6	67.8	-21.2	
14	15	Sampdoria	10	6	22	46	63	-17	36	0.95	37.0	58.5	-21.5	
15	16	Spezia	10	6	22	41	71	-30	36	0.95	39.3	67.6	-28.3	
16	17	Salernitana	7	10	21	33	78	-45	31	0.82	37.7	65.9	-28.2	
17	18	Cagliari	6	12	20	34	68	-34	30	0.79	39.5	61.5	-22.0	
18	19	Genoa	4	16	18	27	60	-33	28	0.74	37.8	51.6	-13.8	
19	20	Venezia	6	9	23	34	69	-35	27	0.71	36.1	72.6	-36.5	


	xGD/90	Attendance
0	0.74	44015
1	1.11	44473
2	0.73	28119
3	0.35	22621
4	0.20	23263
5	0.66	41929
6	0.39	21107
7	0.55	10447
8	0.08	13894
9	0.23	9846
10	-0.29	8362
11	0.03	12144
12	-0.27	14158
13	-0.56	6356
14	-0.57	9417
15	-0.74	6709
16	-0.74	15073
17	-0.58	9718
18	-0.36	12326
19	-0.96	6648

```
[6]: table_df.shape
```


```
[6]: (20, 15)
```

```
[19]: sns.set_style('darkgrid')
matplotlib.rcParams['font.size'] = 14
matplotlib.rcParams['figure.figsize'] = (20, 15)
matplotlib.rcParams['figure.facecolor'] = '#00000000'
```

```
[20]: sns.scatterplot(data=table_df, x='W', y='Rk', hue='Squad', s=50);
```



```
[21]: sns.heatmap(table_df.corr(), annot=True);
```



1.2 Training Preparation

Defining inputs and targets

[22]: `table_df`

```
[22]:
```

	Rk	Squad	W	D	L	GF	GA	GD	Pts	Pts/MP	xG	xGA	xGD	\
0	1	Milan	26	8	4	69	31	38	86	2.26	63.1	34.8	28.3	
1	2	Inter	25	9	4	84	32	52	84	2.21	81.4	39.2	42.2	
2	3	Napoli	24	7	7	74	31	43	79	2.08	59.4	31.9	27.6	
3	4	Juventus	20	10	8	57	37	20	70	1.84	51.5	38.0	13.4	
4	5	Lazio	18	10	10	77	58	19	64	1.68	55.8	48.3	7.6	
5	6	Roma	18	9	11	59	43	16	63	1.66	63.7	38.5	25.2	
6	7	Fiorentina	19	5	14	59	51	8	62	1.63	58.8	44.1	14.7	
7	8	Atalanta	16	11	11	65	48	17	59	1.55	66.2	45.2	20.9	
8	9	Hellas Verona	14	11	13	65	59	6	53	1.39	51.8	48.9	2.9	
9	10	Torino	13	11	14	46	41	5	50	1.32	48.8	39.9	8.9	
10	11	Sassuolo	13	11	14	64	66	-2	50	1.32	56.0	67.0	-10.9	
11	12	Udinese	11	14	13	61	58	3	47	1.24	53.1	52.0	1.0	
12	13	Bologna	12	10	16	44	55	-11	46	1.21	44.1	54.2	-10.1	

13	14	Empoli	10	11	17	50	70	-20	41	1.08	46.6	67.8	-21.2
14	15	Sampdoria	10	6	22	46	63	-17	36	0.95	37.0	58.5	-21.5
15	16	Spezia	10	6	22	41	71	-30	36	0.95	39.3	67.6	-28.3
16	17	Salernitana	7	10	21	33	78	-45	31	0.82	37.7	65.9	-28.2
17	18	Cagliari	6	12	20	34	68	-34	30	0.79	39.5	61.5	-22.0
18	19	Genoa	4	16	18	27	60	-33	28	0.74	37.8	51.6	-13.8
19	20	Venezia	6	9	23	34	69	-35	27	0.71	36.1	72.6	-36.5

	xGD/90	Attendance
0	0.74	44015
1	1.11	44473
2	0.73	28119
3	0.35	22621
4	0.20	23263
5	0.66	41929
6	0.39	21107
7	0.55	10447
8	0.08	13894
9	0.23	9846
10	-0.29	8362
11	0.03	12144
12	-0.27	14158
13	-0.56	6356
14	-0.57	9417
15	-0.74	6709
16	-0.74	15073
17	-0.58	9718
18	-0.36	12326
19	-0.96	6648

```
[27]: input_cols = list(table_df.columns[2:])
      target_col = 'Rk'
```

```
[31]: print('Input:', input_cols)
      print('Target:', target_col)
```

```
Input: ['W', 'D', 'L', 'GF', 'GA', 'GD', 'Pts', 'Pts/MP', 'xG', 'xGA', 'xGD',
       'xGD/90', 'Attendance']
Target: Rk
```

Creating the Two Different Dataframes

```
[43]: inputs_df = table_df[input_cols].copy()
      target_df = table_df[target_col]
```

```
[44]: inputs_df
```

```
[44]:
```

	W	D	L	GF	GA	GD	Pts	Pts/MP	xG	xGA	xGD	xGD/90	Attendance
0	26	8	4	69	31	38	86	2.26	63.1	34.8	28.3	0.74	44015
1	25	9	4	84	32	52	84	2.21	81.4	39.2	42.2	1.11	44473
2	24	7	7	74	31	43	79	2.08	59.4	31.9	27.6	0.73	28119
3	20	10	8	57	37	20	70	1.84	51.5	38.0	13.4	0.35	22621
4	18	10	10	77	58	19	64	1.68	55.8	48.3	7.6	0.20	23263
5	18	9	11	59	43	16	63	1.66	63.7	38.5	25.2	0.66	41929
6	19	5	14	59	51	8	62	1.63	58.8	44.1	14.7	0.39	21107
7	16	11	11	65	48	17	59	1.55	66.2	45.2	20.9	0.55	10447
8	14	11	13	65	59	6	53	1.39	51.8	48.9	2.9	0.08	13894
9	13	11	14	46	41	5	50	1.32	48.8	39.9	8.9	0.23	9846
10	13	11	14	64	66	-2	50	1.32	56.0	67.0	-10.9	-0.29	8362
11	11	14	13	61	58	3	47	1.24	53.1	52.0	1.0	0.03	12144
12	12	10	16	44	55	-11	46	1.21	44.1	54.2	-10.1	-0.27	14158
13	10	11	17	50	70	-20	41	1.08	46.6	67.8	-21.2	-0.56	6356
14	10	6	22	46	63	-17	36	0.95	37.0	58.5	-21.5	-0.57	9417
15	10	6	22	41	71	-30	36	0.95	39.3	67.6	-28.3	-0.74	6709
16	7	10	21	33	78	-45	31	0.82	37.7	65.9	-28.2	-0.74	15073
17	6	12	20	34	68	-34	30	0.79	39.5	61.5	-22.0	-0.58	9718
18	4	16	18	27	60	-33	28	0.74	37.8	51.6	-13.8	-0.36	12326
19	6	9	23	34	69	-35	27	0.71	36.1	72.6	-36.5	-0.96	6648

```
[37]: target_df
```

```
[37]: 0      1
      1      2
      2      3
      3      4
      4      5
      5      6
      6      7
      7      8
      8      9
      9     10
     10     11
     11     12
     12     13
     13     14
     14     15
     15     16
     16     17
     17     18
     18     19
     19     20
Name: Rk, dtype: int64
```

1.2.1 Scaling Numerical Values

```
[46]: numerical_cols = list(inputs_df.columns)
```

```
[47]: numerical_cols
```

```
[47]: ['W',  
      'D',  
      'L',  
      'GF',  
      'GA',  
      'GD',  
      'Pts',  
      'Pts/MP',  
      'xG',  
      'xGA',  
      'xGD',  
      'xGD/90',  
      'Attendance']
```

```
[38]: from sklearn.preprocessing import MinMaxScaler
```

```
[49]: scaler = MinMaxScaler().fit(inputs_df[numerical_cols])
```

```
[50]: inputs_df[numerical_cols] = scaler.transform(inputs_df[numerical_cols])
```

```
[51]: inputs_df
```

```
[51]:
```

	W	D	L	GF	GA	GD	Pts	\
0	1.000000	0.272727	0.000000	0.736842	0.000000	0.855670	1.000000	
1	0.954545	0.363636	0.000000	1.000000	0.021277	1.000000	0.966102	
2	0.909091	0.181818	0.157895	0.824561	0.000000	0.907216	0.881356	
3	0.727273	0.454545	0.210526	0.526316	0.127660	0.670103	0.728814	
4	0.636364	0.454545	0.315789	0.877193	0.574468	0.659794	0.627119	
5	0.636364	0.363636	0.368421	0.561404	0.255319	0.628866	0.610169	
6	0.681818	0.000000	0.526316	0.561404	0.425532	0.546392	0.593220	
7	0.545455	0.545455	0.368421	0.666667	0.361702	0.639175	0.542373	
8	0.454545	0.545455	0.473684	0.666667	0.595745	0.525773	0.440678	
9	0.409091	0.545455	0.526316	0.333333	0.212766	0.515464	0.389831	
10	0.409091	0.545455	0.526316	0.649123	0.744681	0.443299	0.389831	
11	0.318182	0.818182	0.473684	0.596491	0.574468	0.494845	0.338983	
12	0.363636	0.454545	0.631579	0.298246	0.510638	0.350515	0.322034	
13	0.272727	0.545455	0.684211	0.403509	0.829787	0.257732	0.237288	
14	0.272727	0.090909	0.947368	0.333333	0.680851	0.288660	0.152542	
15	0.272727	0.090909	0.947368	0.245614	0.851064	0.154639	0.152542	
16	0.136364	0.454545	0.894737	0.105263	1.000000	0.000000	0.067797	
17	0.090909	0.636364	0.842105	0.122807	0.787234	0.113402	0.050847	
18	0.000000	1.000000	0.736842	0.000000	0.617021	0.123711	0.016949	


```
19 0.090909 0.363636 1.000000 0.122807 0.808511 0.103093 0.000000
```

	Pts/MP	xG	xGA	xGD	xGD/90	Attendance
0	1.000000	0.596026	0.071253	0.823380	0.821256	0.987984
1	0.967742	1.000000	0.179361	1.000000	1.000000	1.000000
2	0.883871	0.514349	0.000000	0.814485	0.816425	0.570953
3	0.729032	0.339956	0.149877	0.634053	0.632850	0.426712
4	0.625806	0.434879	0.402948	0.560356	0.560386	0.443555
5	0.612903	0.609272	0.162162	0.783990	0.782609	0.933258
6	0.593548	0.501104	0.299754	0.650572	0.652174	0.386993
7	0.541935	0.664459	0.326781	0.729352	0.729469	0.107327
8	0.438710	0.346578	0.417690	0.500635	0.502415	0.197760
9	0.393548	0.280353	0.196560	0.576874	0.574879	0.091560
10	0.393548	0.439294	0.862408	0.325286	0.323671	0.052627
11	0.341935	0.375276	0.493857	0.476493	0.478261	0.151848
12	0.322581	0.176600	0.547912	0.335451	0.333333	0.204686
13	0.238710	0.231788	0.882064	0.194409	0.193237	0.000000
14	0.154839	0.019868	0.653563	0.190597	0.188406	0.080305
15	0.154839	0.070640	0.877150	0.104193	0.106280	0.009261
16	0.070968	0.035320	0.835381	0.105464	0.106280	0.228691
17	0.051613	0.075055	0.727273	0.184244	0.183575	0.088202
18	0.019355	0.037528	0.484029	0.288437	0.289855	0.156623
19	0.000000	0.000000	1.000000	0.000000	0.000000	0.007661

1.3 Training, Validation and Test Set

```
[52]: from sklearn.model_selection import train_test_split
```

```
[57]: train_inputs, val_inputs, train_targets, val_targets = \
      ↪train_test_split(inputs_df, target_df, test_size=0.20, random_state=42)
```

```
[58]: train_inputs
```

```
[58]:
```

	W	D	L	GF	GA	GD	Pts \
8	0.454545	0.545455	0.473684	0.666667	0.595745	0.525773	0.440678
5	0.636364	0.363636	0.368421	0.561404	0.255319	0.628866	0.610169
11	0.318182	0.818182	0.473684	0.596491	0.574468	0.494845	0.338983
3	0.727273	0.454545	0.210526	0.526316	0.127660	0.670103	0.728814
18	0.000000	1.000000	0.736842	0.000000	0.617021	0.123711	0.016949
16	0.136364	0.454545	0.894737	0.105263	1.000000	0.000000	0.067797
13	0.272727	0.545455	0.684211	0.403509	0.829787	0.257732	0.237288
2	0.909091	0.181818	0.157895	0.824561	0.000000	0.907216	0.881356
9	0.409091	0.545455	0.526316	0.333333	0.212766	0.515464	0.389831
19	0.090909	0.363636	1.000000	0.122807	0.808511	0.103093	0.000000
4	0.636364	0.454545	0.315789	0.877193	0.574468	0.659794	0.627119
12	0.363636	0.454545	0.631579	0.298246	0.510638	0.350515	0.322034
7	0.545455	0.545455	0.368421	0.666667	0.361702	0.639175	0.542373

10	0.409091	0.545455	0.526316	0.649123	0.744681	0.443299	0.389831
14	0.272727	0.090909	0.947368	0.333333	0.680851	0.288660	0.152542
6	0.681818	0.000000	0.526316	0.561404	0.425532	0.546392	0.593220

	Pts/MP	xG	xGA	xGD	xGD/90	Attendance
8	0.438710	0.346578	0.417690	0.500635	0.502415	0.197760
5	0.612903	0.609272	0.162162	0.783990	0.782609	0.933258
11	0.341935	0.375276	0.493857	0.476493	0.478261	0.151848
3	0.729032	0.339956	0.149877	0.634053	0.632850	0.426712
18	0.019355	0.037528	0.484029	0.288437	0.289855	0.156623
16	0.070968	0.035320	0.835381	0.105464	0.106280	0.228691
13	0.238710	0.231788	0.882064	0.194409	0.193237	0.000000
2	0.883871	0.514349	0.000000	0.814485	0.816425	0.570953
9	0.393548	0.280353	0.196560	0.576874	0.574879	0.091560
19	0.000000	0.000000	1.000000	0.000000	0.000000	0.007661
4	0.625806	0.434879	0.402948	0.560356	0.560386	0.443555
12	0.322581	0.176600	0.547912	0.335451	0.333333	0.204686
7	0.541935	0.664459	0.326781	0.729352	0.729469	0.107327
10	0.393548	0.439294	0.862408	0.325286	0.323671	0.052627
14	0.154839	0.019868	0.653563	0.190597	0.188406	0.080305
6	0.593548	0.501104	0.299754	0.650572	0.652174	0.386993

[59]: train_targets

```
[59]: 8      9
      5      6
      11     12
      3      4
      18     19
      16     17
      13     14
      2      3
      9     10
      19     20
      4      5
      12     13
      7      8
      10     11
      14     15
      6      7
      Name: Rk, dtype: int64
```

[60]: val_inputs

```
[60]:           W           D           L           GF           GA           GD           Pts  \
0      1.000000  0.272727  0.000000  0.736842  0.000000  0.855670  1.000000
17  0.090909  0.636364  0.842105  0.122807  0.787234  0.113402  0.050847
```

```

15  0.272727  0.090909  0.947368  0.245614  0.851064  0.154639  0.152542
1   0.954545  0.363636  0.000000  1.000000  0.021277  1.000000  0.966102

```

```

      Pts/MP      xG      xGA      xGD      xGD/90  Attendance
0   1.000000  0.596026  0.071253  0.823380  0.821256    0.987984
17  0.051613  0.075055  0.727273  0.184244  0.183575    0.088202
15  0.154839  0.070640  0.877150  0.104193  0.106280    0.009261
1   0.967742  1.000000  0.179361  1.000000  1.000000    1.000000

```

```
[61]: val_targets
```

```

[61]: 0      1
      17     18
      15     16
      1      2
      Name: Rk, dtype: int64

```

2 Train a Linear Regression Model

```

[78]: from sklearn.linear_model import Ridge
      from sklearn.metrics import mean_squared_error

```

2.0.1 Train

```

[79]: %%time
      model = Ridge().fit(train_inputs, train_targets)

```

```

CPU times: total: 0 ns
Wall time: 14.1 ms

```

```
[80]: train_preds = model.predict(train_inputs)
```

```
[81]: train_preds
```

```

[81]: array([ 9.77093948,  5.88979012, 11.15823153,  6.04573319, 16.92544041,
           17.26138736, 14.43950518,  2.3812202 , 10.38990622, 18.16848165,
           6.86681754, 12.57885414,  7.52433463, 11.52068121, 14.79206149,
           7.28661564])

```

```
[82]: train_targets
```

```

[82]: 8      9
      5      6
      11     12
      3      4
      18     19
      16     17

```

```

13    14
2      3
9     10
19    20
4      5
12    13
7      8
10    11
14    15
6      7
Name: Rk, dtype: int64

```

```
[83]: train_rmse = mean_squared_error(train_preds, train_targets, squared=False)
print('The position is wrong of {} places'.format(train_rmse))
```

The position is wrong of 1.0677042461643216 places

2.0.2 Validation

```
[84]: model = Ridge().fit(val_inputs, val_targets)
```

```
[85]: val_preds = model.predict(val_inputs)
```

```
[86]: val_preds
```

```
[86]: array([ 2.86574459, 16.28920059, 15.87623286,  1.96882195])
```

```
[87]: val_targets
```

```

[87]: 0      1
17     18
15     16
1      2
Name: Rk, dtype: int64

```

```
[88]: val_rmse = mean_squared_error(val_preds, val_targets, squared=False)
print('The position is wrong of {} places'.format(val_rmse))
```

The position is wrong of 1.2672931673486385 places

```
[89]: val_rmse
```

```
[89]: 1.2672931673486385
```

2.0.3 Evaluating the Weights

```
[102]: weights = model.coef_
```

```
[110]: weights_df = pd.DataFrame({'Parameters': train_inputs.columns,
                                'Weights': weights})
```

```
[115]: weights_df.sort_values('Weights', ascending=False)
```

```
[115]:
```

	Parameters	Weights
2	L	1.557926
4	GA	1.452797
9	xGA	1.225432
1	D	0.602440
8	xG	-0.943708
3	GF	-1.103084
11	xGD/90	-1.173862
10	xGD	-1.176938
5	GD	-1.352136
0	W	-1.646702
12	Attendance	-1.660106
7	Pts/MP	-1.728109
6	Pts	-1.729754

These are the different weights' parameters.

3 Making Predictions

```
[90]: def make_preds(user_input):
        input_df = pd.DataFrame([user_input])
        input_df[numerical_cols] = scaler.transform(input_df[numerical_cols])
        return model.predict(input_df[numerical_cols])
```

Let's see if it can properly guess in what position Roma finished its 2018/2019 season.

```
[97]: user_input = {'W':18,
                    'D':12, 'L':8, 'GF':66, 'GA':48, 'GD': 18,
                    'Pts':66, 'Pts/MP':1.74, 'xG':64.3, 'xGA':54.4, 'xGD':9.9,
                    'xGD/90':0.26, 'Attendance': 38622}
```

```
[98]: make_preds(user_input)
```

```
[98]: array([7.38742012])
```

Quite, good, Roma final position was 6.